BookMyShow Scalable Platform – Project Report

Objective:

To simulate a BookMyShow-like ticketing system and test performance under high traffic, identifying bottlenecks, optimizing architecture, and ensuring scalability, resilience, and cost efficiency.

Scope:

- Create a prototype with frontend and backend
- Deploy backend on Cloud Run
- Use Cloud SQL for database
- Perform load testing with k6
- Implement caching and CDN
- Monitor performance and analyze costs

System Architecture

Components:

- Frontend: HTML/CSS/JS hosted in Cloud Storage (CDN enabled)
- Backend: Node.js + Express, deployed on Cloud Run
- Database: PostgreSQL (Cloud SQL)
- CDN: Cloud CDN on a backend bucket for static content
- Cloud Load Balancing: Handled by Cloud Run autoscaling
- Monitoring: Cloud Monitoring and Logging

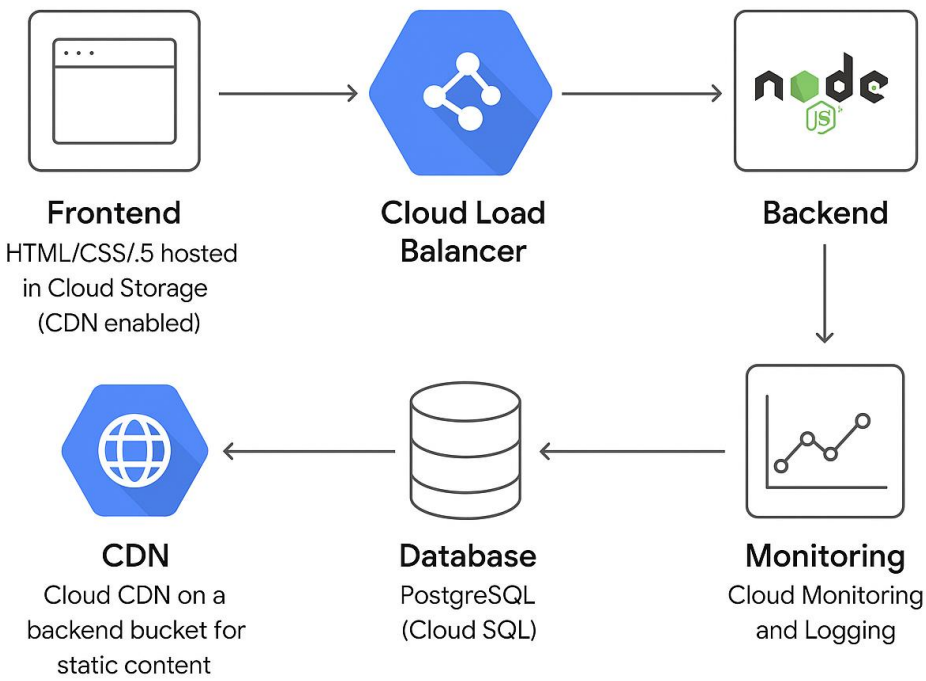

Deployment & Code

Backend (server.js):

Frontend: HTML, CSS, JS hosted on Cloud Storage + served via Cloud CDN.

```
const express = require("express");
const path = require("path");
const { Pool } = require("pg");
const app = express();
const PORT = process.env.PORT || 3000;

// Serve static files
app.use(express.static(path.join(__dirname, "public")));

// Database configuration
const pool = new Pool({
  user: process.env.DB_USER,
  host: process.env.DB_HOST,
  database: process.env.DB_NAME,
  password: process.env.DB_PASS,
  port: 5432,
});

// Test DB connection
pool.connect((err, client, release) => {
  if (err) console.error("Error connecting to DB:", err.stack);
  else console.log("Connected to PostgreSQL DB");
  release();
});

// Example API route
app.get("/api/movies", async (req, res) => {
  try {
    const result = await pool.query("SELECT * FROM movies");
    res.json(result.rows);
  } catch (err) {
    console.error(err);
    res.status(500).json({ error: "Database error" });
  }
});

// Fallback for index.html
app.get("/", (req, res) => {
  res.sendFile(path.join(__dirname, "public", "index.html"));
});

app.listen(PORT, () => console.log(`Server running on port ${PORT}`));
```

Docker:

- Image built locally → pushed to Google Container Registry
- Deployed to Cloud Run with Cloud SQL integration

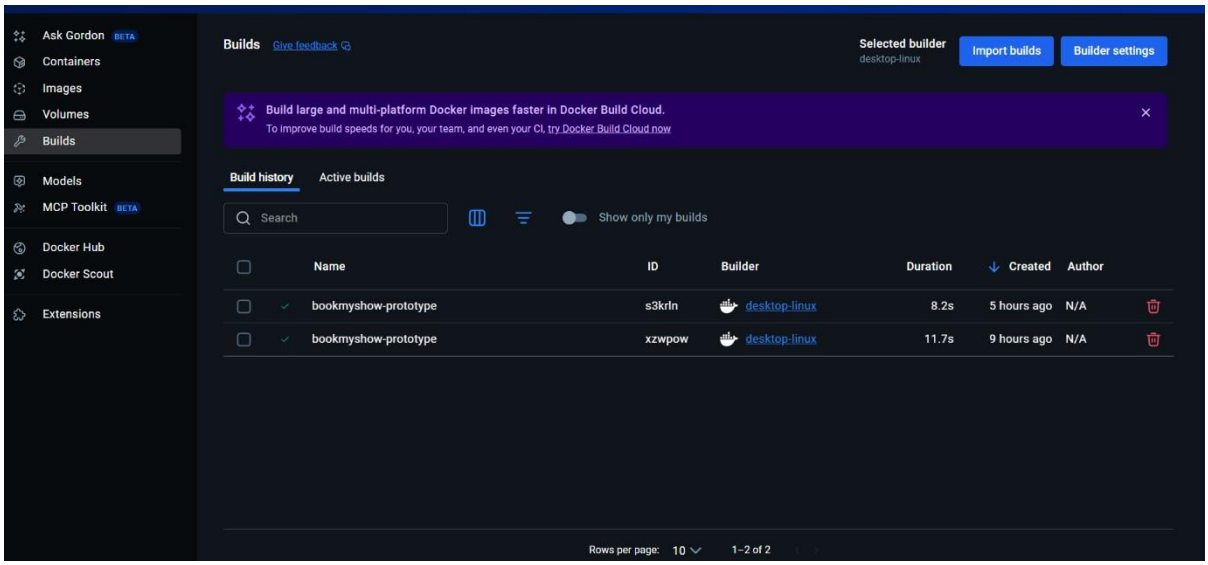

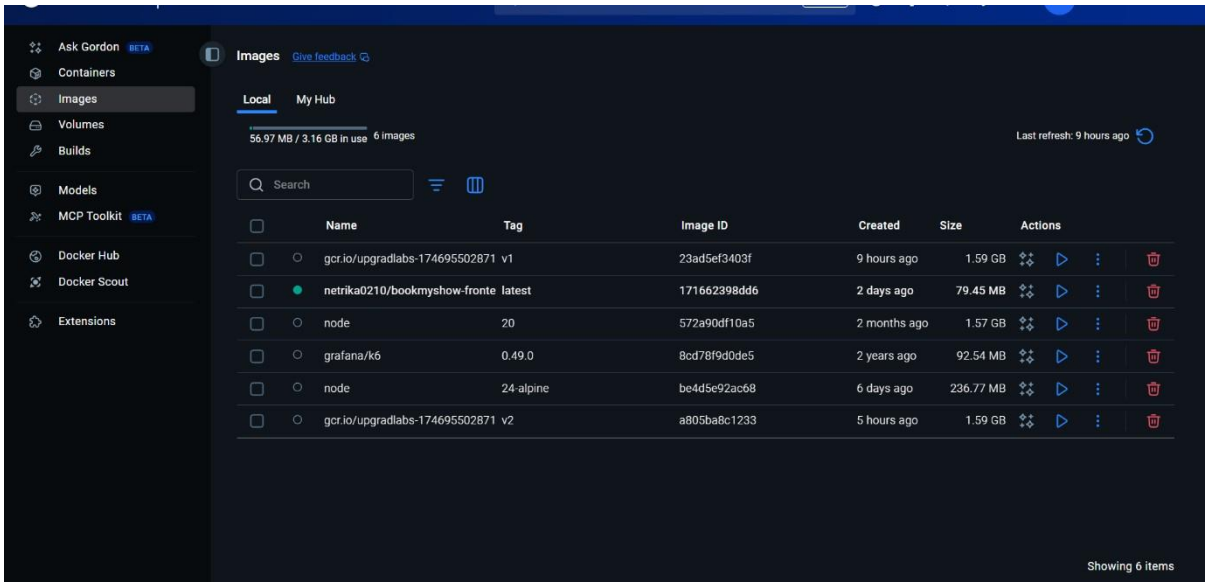

Traffic Simulation

Tool: k6 load testing

Metrics for v2 Deployment:

Metrics	Value
Avg Response Time	766ms
Max Response Time	2.37s
Iterations	1748
VUs	50

Observations:

- Cold starts on Cloud Run may cause initial delays
- Database response time ~765ms, potential for query optimization

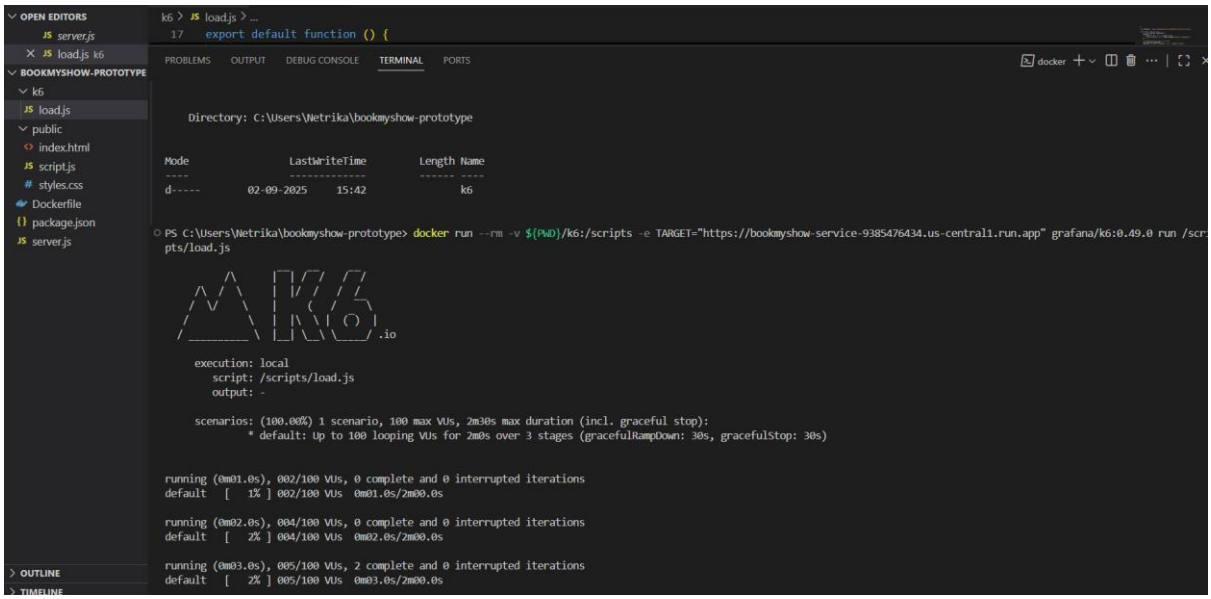

Database Optimisation

Actions Taken / Recommendations:

- Cloud SQL instance with PostgreSQL 15 • Database connection pooling via pg module
- **Future enhancements:**
  - Horizontal read replicas for high availability
  - Backups and point-in-time recovery ○ Query indexing for faster reads

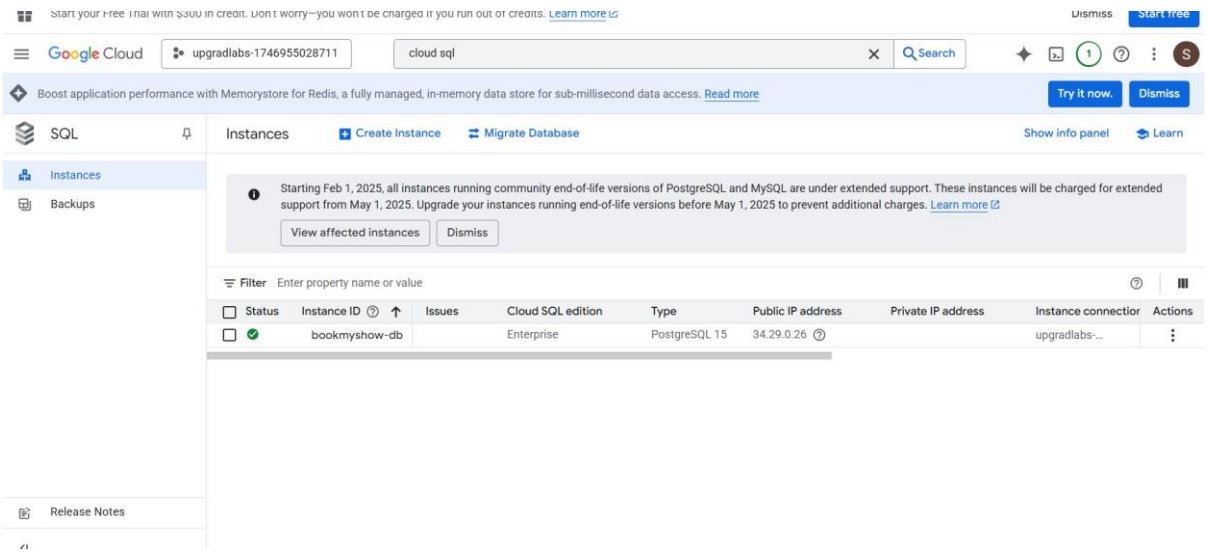

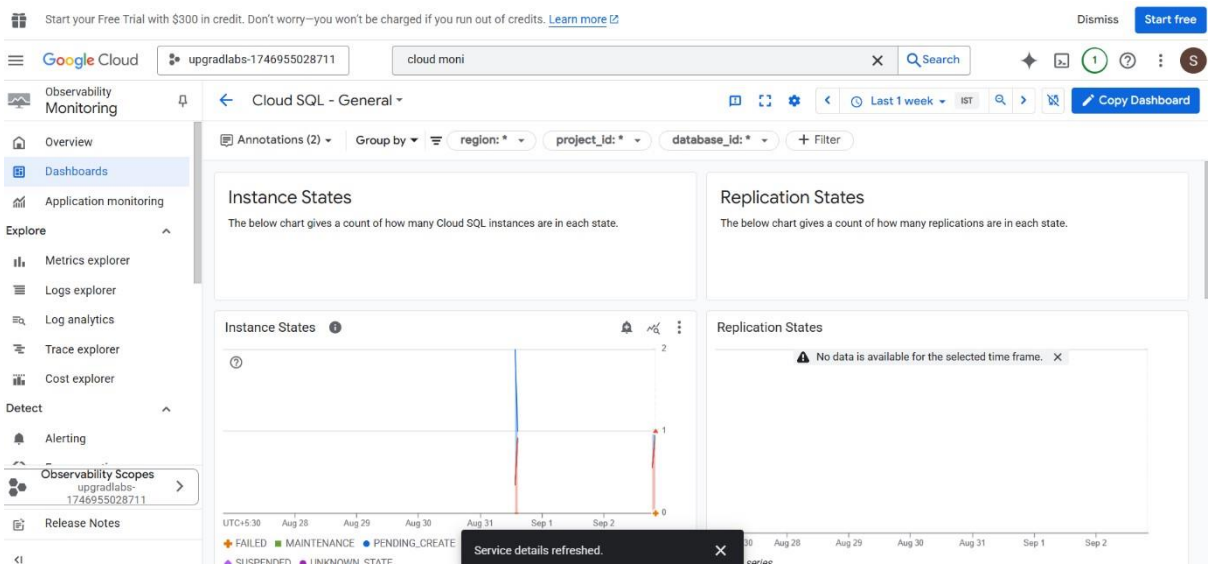

Application Resilience

- Static content cached in Cloud Storage + Cloud CDN • Autoscaling enabled on Cloud Run
- Future enhancements:
  - Rate limiting
  - Circuit breakers for backend APIs ○ Multi-region deployment

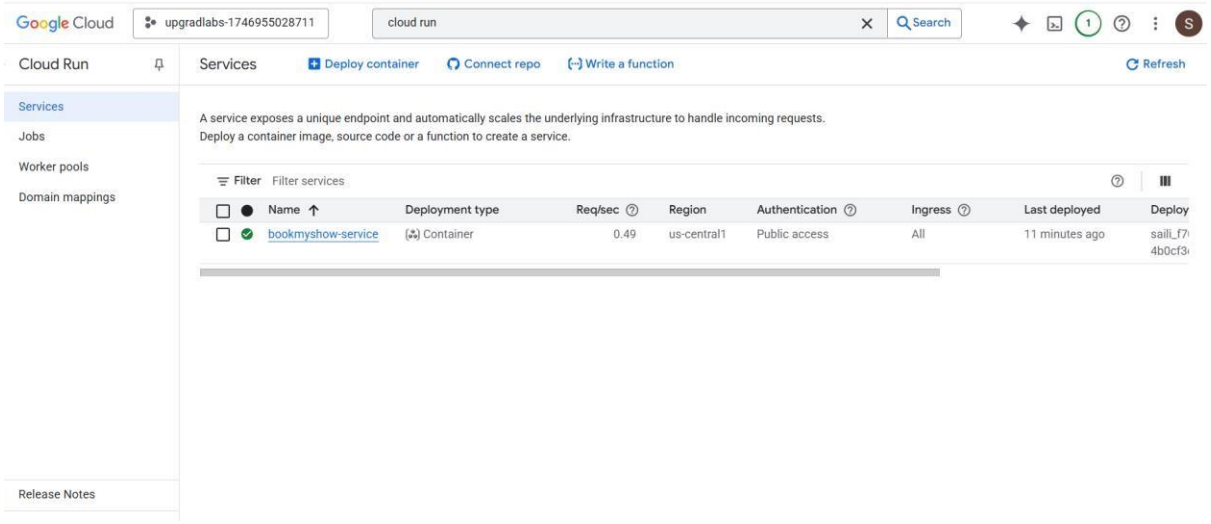
Google Cloud
upgradlabs-1746955028711
cloud run
Search
Cloud Run
Services
Deploy container
Connect repo
Write a function
Refresh
Services
Jobs
Worker pools
Domain mappings
A service exposes a unique endpoint and automatically scales the underlying infrastructure to handle incoming requests.
Deploy a container image, source code or a function to create a service.
Filter
Filter services
Name
Deployment type
Req/sec
Region
Authentication
Ingress
Last deployed
bookmyshow-service
Container
0.49
us-central1
Public access
All
11 minutes ago
Release Notes

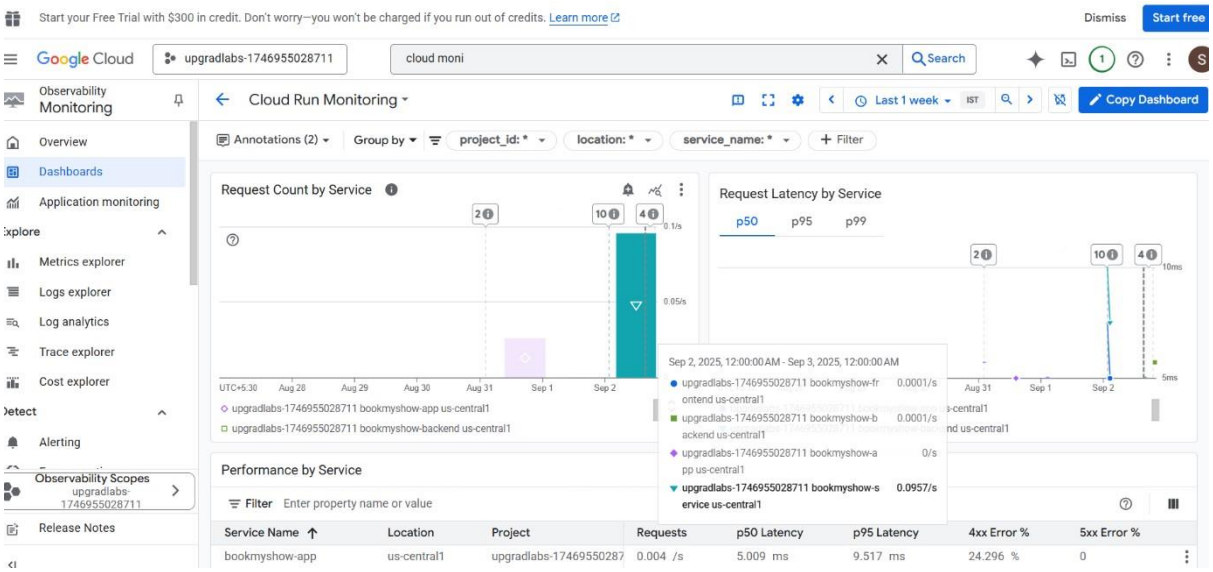
Start your Free Trial with $300 in credit. Don't worry—you won't be charged if you run out of credits. Learn more
Dismiss
Start free
Google Cloud
upgradlabs-1746955028711
cloud moni
Search
Observability
Monitoring
Cloud Run Monitoring
Last 1 week
Copy Dashboard
Overview
Dashboards
Application monitoring
Metrics explorer
Logs explorer
Log analytics
Trace explorer
Cost explorer
Alerting
Observability Scopes
upgradlabs-1746955028711
Release Notes
Annotations (2)
Group by
project_id: *
location: *
service_name: *
Filter
Request Count by Service
Request Latency by Service
p50
p95
p99
Performance by Service
Filter
Enter property name or value
Service Name
Location
Project
Requests
p50 Latency
p95 Latency
4xx Error %
5xx Error %
bookmyshow-app
us-central1
upgradlabs-17469550287
0.004 /s
5.009 ms
9.517 ms
24.296 %
0

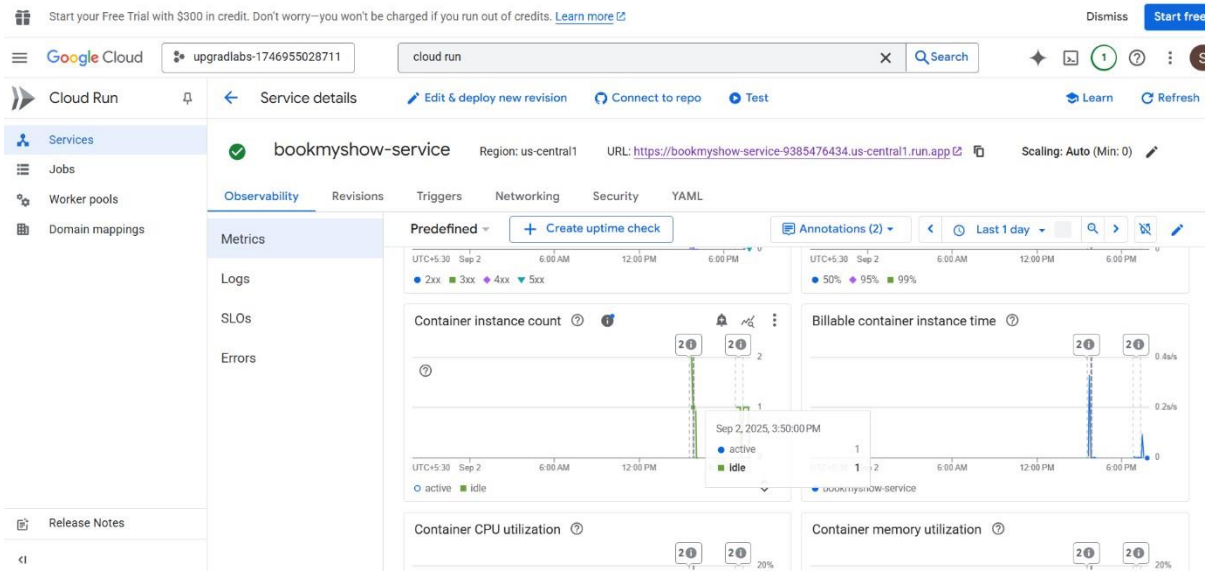
Start your Free Trial with $300 in credit. Don't worry—you won't be charged if you run out of credits. Learn more
Dismiss
Start free
Google Cloud
upgradlabs-1746955028711
cloud run
Search
Cloud Run
Service details
Edit & deploy new revision
Connect to repo
Test
Learn
Refresh
Services
Jobs
Worker pools
Domain mappings
bookmyshow-service
Region: us-central1
URL: https://bookmyshow-service-9385476434.us-central1.run.app
Scaling: Auto (Min: 0)
Observability
Revisions
Triggers
Networking
Security
YAML
Metrics
Logs
SLOs
Errors
Predefined
Create uptime check
Annotations (2)
Last 1 day
Container instance count
Billable container instance time
Container CPU utilization
Container memory utilization
Release Notes

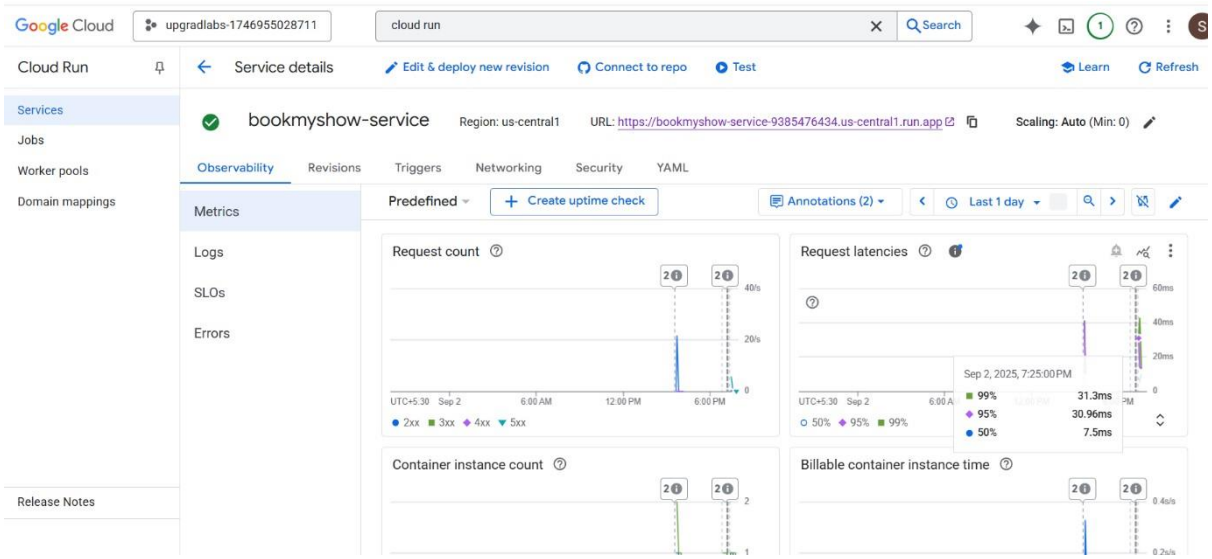

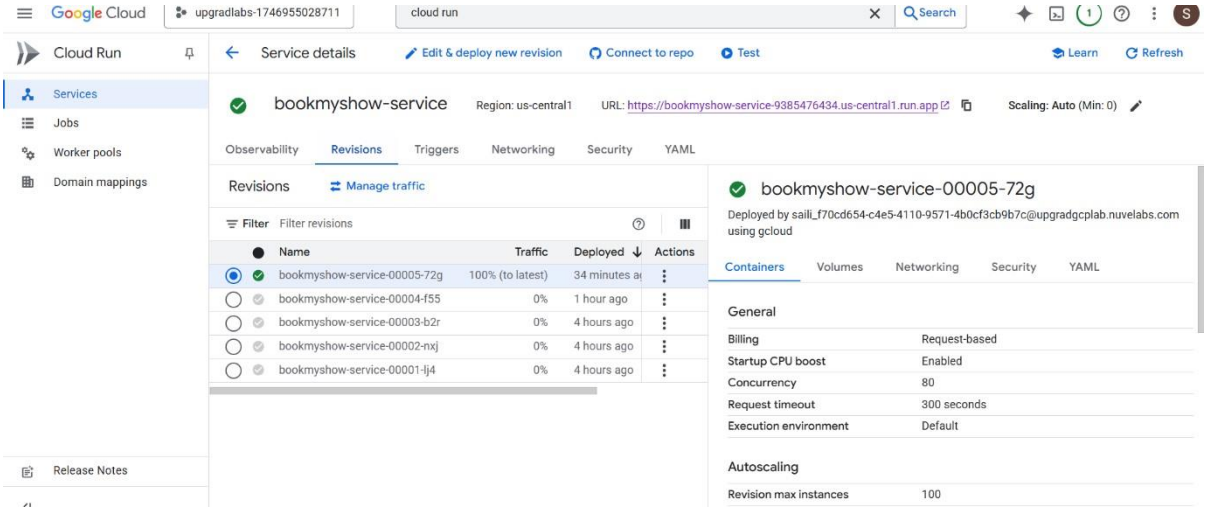

Monitoring & Alerts

Metrics monitored: CPU, Memory, HTTP request latency

- Cloud Monitoring dashboard
- Alert configuration for thresholds (e.g., CPU > 80%)

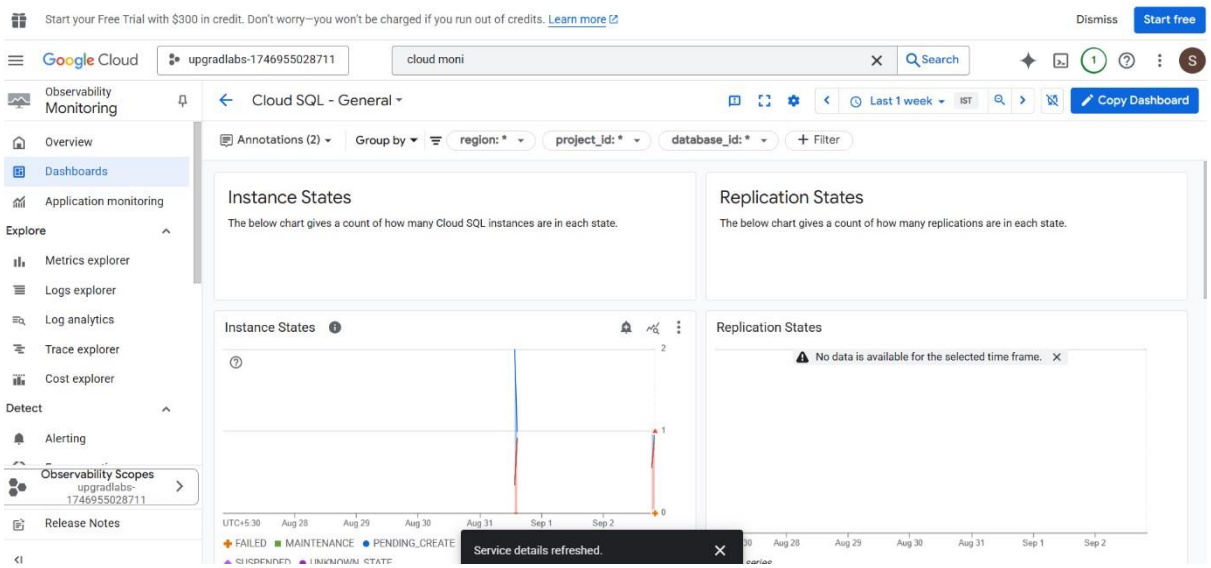

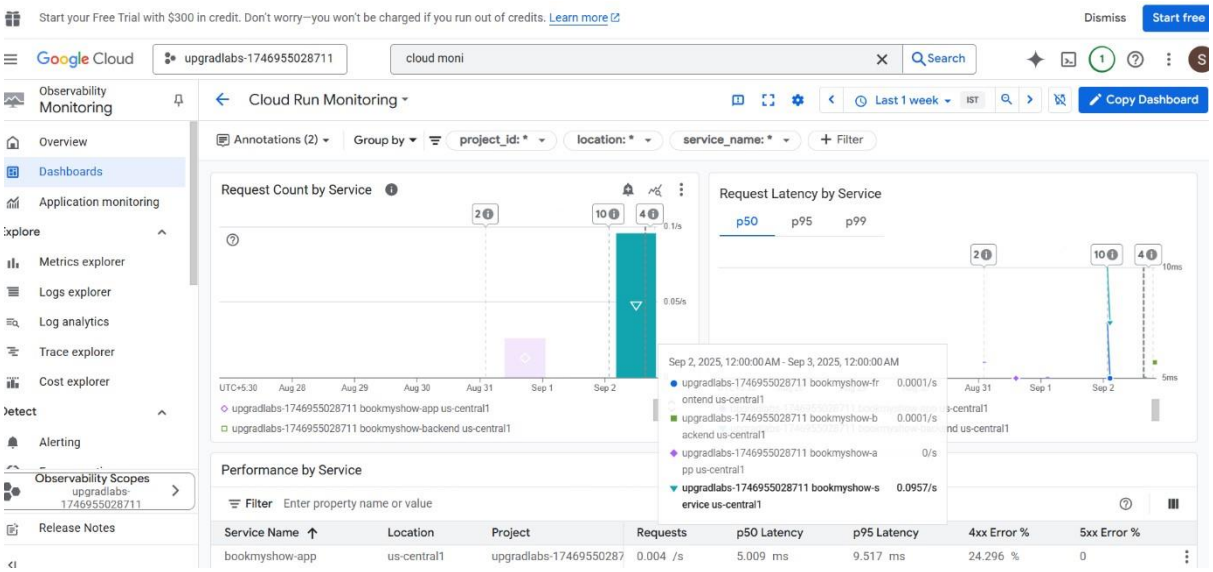

Cost & Budget Analysis **Free**

Tier Usage:

Service	Usage	Estimated Cost
Cloud Run	1 vCPU, 512MB	$0
Cloud SQL	db-f1-micro	$0
Storage & CDN	3 files, 2.9 KiB	$0

Budget Alerts:

- Budget: $10/month
- Thresholds: 50%, 90%, 100%
- Email notifications enabled (mockup screenshot) **Cost optimization tips:**
- Use preemptible VMs for batch tasks

- Leverage Cloud Run scaling efficiently
- Free tier ensures minimal cost

Final Evaluation

Performance Improvements:

- v2 backend deployed on Cloud Run shows reduced cold start impact
- Static content served from CDN improves load times **Architecture:**
- Cloud Run + Cloud SQL + Cloud Storage provides scalable solution

Future Recommendations:

- Implement read replicas and query optimization
- Set up real monitoring and alerting
- Introduce chaos testing for resilience
- Evaluate cost savings via sustained use discounts

GitHub Repository

Repo: https://github.com/netrikadongre-source/bookmyshowproject

Contains:

- **Full backend code (server.js, package.json, Dockerfile)**
- **Frontend (index.html, styles.css, script.js)**
- **Instructions for deployment and testing**

Author: Netrika Dongre